

# SQL Interview Prep Questions - 2 (16 - 30)

16) Write an SQL query to find the people who have the most friends and the most friends number.

Difficulty: Medium

Company : Facebook

Table: RequestAccepted

Column Name	Type
requester_id	int
accepter_id	int
accept_date	date

(requester\_id, accepter\_id) is the primary key for this table.

This table contains the ID of the user who sent the request, the ID of the user who received the request, and the date when the request was accepted.

**Input:**

RequestAccepted table:

requester_id	accepter_id	accept_date
1	2	2016/06/03
1	3	2016/06/08
2	3	2016/06/08
3	4	2016/06/09

**Output:**

id	num
3	3

**Explanation:**

A user can be both a requester as well as an acceptor. So the friend count for a user should be accounted for when the user is on the requestor side or the acceptor side.

**Step1:** Find the count of friends for each user when they are on the requestor side

**Step2:** Similarly find the count of friends when the users are on the acceptor side

**Step3:** Combine the count of users' friends when they were a requestor / acceptor by taking a union of both the tables.

**Step4:** Find the total count by taking the sum of friends count by taking a group by on each users

**Step5:** Find the user with the maximum number of friends using the filters.

```
with requestor as (  
  select requester_id as user_id, count(accepter_id) frnd_count from RequestAccepted  
  group by requester_id  
,  
  
  acceptor as (  
    select accepter_id as user_id, count(requester_id) frnd_count from RequestAccepted  
    group by accepter_id  
  ),  
  
  all_frnd_count as  
  (  
    select * from requestor union all select * from acceptor  
  ),  
  
  uniq_user_frnd_cnt as (  
    select user_id as id, sum(frnd_count) as num from all_frnd_count group by user_id  
  )  
  
select * from uniq_user_frnd_cnt  
where num = (select max(num) from uniq_user_frnd_cnt)
```

---

17) A quiet student is the one who took at least one exam and did not score the high or the low score.

Write an SQL query to report the students (student\_id, student\_name) being quiet in all exams. Do not return the student who has never taken any exam.

Return the result table ordered by student\_id.

Difficulty: Hard

Company: Yahoo

Table: Student

Column Name	Type
student_id	int
student_name	varchar

student\_id is the primary key for this table.  
student\_name is the name of the student.

Table: Exam

Column Name	Type
exam_id	int
student_id	int
score	int

(exam\_id, student\_id) is the primary key for this table.

Each row of this table indicates that the student with student\_id had a score points in the exam with id exam\_id.

### **Explanation:**

**Step1:** Find all the students whose marks are the highest and lowest in any of the exam so they can be eliminated. This can be found by ranking the marks using the windows function in the exam table.

**Step2:** For finding the students who have scored the lowest marks partition by exam\_id and sort marks in the ascending order so the students with the lowest mark are ranked 1.

**Step3:** Similarly create another column by partitioning on the exam\_id and in the descending order of the marks so students with the highest marks will be ranked 1.

**Step4:** For a student to be silent there are two conditions given below which can be fed as subqueries in the filter condition.

1. Should have taken at least one test i.e. student\_id should be in the exam table
2. Should not be highest or lowest scorer in any exam i.e. Should not be ranked 1 in step 2 or step 3.

```
Select student_id,  
       Student_name  
  from student  
 where student_id in (select student_id from exam) and student_id not in ( select  
 student_id from( select student_id,  
                      rank() over(partition by exam_id order by score asc) as asc_rank,  
                      rank() over(partition by exam_id order by score desc) as  
                      desc_rank from exam
```

```

) tab1
where asc_rank=1 or desc_rank=1
)

```

---

18) Write an SQL query to report the sum of all total investment values in 2016 tiv\_2016, for all policyholders who:

- have the same tiv\_2015 value as one or more other policyholders, and
- are not located in the same city like any other policyholder (i.e., the (lat, lon) attribute pairs must be unique).

Round tiv\_2016 to two decimal places.

Difficulty: Medium

Company : Twitter

Table: Insurance

Column Name	Type
pid	int
tiv_2015	float
tiv_2016	float
lat	float
lon	float

pid is the primary key column for this table.

Each row of this table contains information about one policy where:

pid is the policyholder's policy ID.

tiv\_2015 is the total investment value in 2015 and tiv\_2016 is the total investment value in 2016.

lat is the latitude of the policy holder's city.

lon is the longitude of the policy holder's city.

**Input:**

Insurance table:

pid	tiv_2015	tiv_2016	lat	lon
1	10	5	10	10
2	20	20	20	20
3	10	30	20	20
4	10	40	40	40

**Output:**

tiv_2016
45.00

**Explanation:**

The first record in the table, like the last record, meets both of the two criteria.

The tiv\_2015 value 10 is the same as the third and fourth records, and its location is unique.

The second record does not meet any of the two criteria. Its tiv\_2015 is not like any other policyholders and its location is the same as the third record, which makes the third record fail, too. So, the result is the sum of tiv\_2016 of the first and last record, which is 45.

**Explanation:**

**Step1:** For the first condition the investment value in 2015 should be the same i.e. it has to be duplicated or in other words unique investment values should not be considered. This can be done by taking a count of the tiv\_2015 column and considering only those that have a count >1.

**Step2:** For the second condition every location should be unique, which means if location occurs more than once in the table, we should not consider it. This can be done by taking a count of the combination of lat and lon and considering only the count =1

**Step3:** Find the overall sum of tiv\_2016 for records matching step1 and step2.

WITH

firstConditionCTE AS

```
(
  SELECT tiv_2015 FROM Insurance GROUP BY tiv_2015
  HAVING COUNT(*) > 1
),
```

SecondConditionCTE AS

```
(
  SELECT concat(lat, ', ', lon) as uniq_lat_lon FROM Insurance
  GROUP BY concat(lat, ', ', lon) HAVING COUNT(*) = 1
)
```

SELECT

```
ROUND(SUM(tiv_2016),2) AS tiv_2016
```

```

FROM Insurance
WHERE tiv_2015 IN (SELECT * FROM firstConditionCTE)
AND
concat(lat, ' ', lon) IN (SELECT uniq_lat_lon FROM SecondConditionCTE)

```

---

19) Write an SQL query to report all possible axis-aligned rectangles with a non-zero area that can be formed by any two points from the Points table.

Each row in the result should contain three columns (p1, p2, area) where:

- p1 and p2 are the id's of the two points that determine the opposite corners of a rectangle.
- area is the area of the rectangle and must be non-zero.

Return the result table ordered by area in descending order. If there is a tie, order them by p1 in ascending order. If there is still a tie, order them by p2 in ascending order.

Difficulty: Medium

Company: Twitter

Table: Points

Column Name	Type
id	int
x_value	int
y_value	int

id is the primary key for this table.

Each point is represented as a 2D coordinate (x\_value, y\_value).

### **Explanation:**


To find the area of a rectangle, we need to know the length and breadth i.e. 2 sides of the rectangle. For this we need 4 points(2 Xs and 2 Ys) to identify the corners of the rectangle. Each row in the Table contains one combination of x and y. So we will combine each row with other rows to get the 4 points.

**Step1:** Since we wanted to find all the possible rectangles with the points in the table, We will do a self join of the table. We use the < expression to join the two tables as we want to find the area for one combination of rectangle only once. I.e. Area of (p1, p2) is the same as (p2, p1).

**Step2:** Let's consider example points  $X_1=4$ ,  $X_2=6$ , and  $X_1 = 4$ ,  $X_2=1$ . As you can see in the first case  $X_1$  is smaller than  $X_2$  and in the second case  $X_1$  is greater than  $X_2$ . Since we want only to find the distance between the two points we take the absolute difference between the two. The same goes with  $Y$  as well.

**Step3:** To avoid getting Area as 0, we filter out points with equal  $X$  values and equal  $Y$  values.

**Example 1:**



```
select
  l.id as P1,
  r.id as P2,
  abs(l.x_value-r.x_value) * abs(l.y_value-r.y_value) as AREA
from points l
join points r on l.id < r.id
```

```
where l.x_value <> r.x_value
and l.y_value <> r.y_value
ORDER BY AREA DESC, p1 ASC, p2 ASC;
```

---

20) Assume you are given the tables below about Facebook pages and page likes. Write a query to return the page IDs of all the Facebook pages that don't have any likes. The output should be in ascending order.

Difficulty: Easy

Company: Facebook

pages Table:	
Column Name	Type
page_id	integer
name	varchar

page_likes Table:	
Column Name	Type
user_id	integer
page_id	integer
liked_date	datetime

**Explanation:**

**Step1:** Since all the pages that have at least one like will be present in the page\_likes Tables. We can write a subquery to get all those pages.

**Step2:** Write a query to find all the pages from the pages table which are not part of the page\_likes.

```
SELECT page_id FROM pages
where page_id not in (select distinct page_id from page_likes);
```

---



21) Assume you have an events table on app analytics. Write a query to get the app's click-through rate (CTR %) in 2022. Output the results in percentages rounded to 2 decimal places.

Notes:

- To avoid integer division, you should multiply the click-through rate by 100.0, not 100.
- Percentage of click-through rate =  $100.0 * \text{Number of clicks} / \text{Number of impressions}$

Company: Facebook

Difficulty: Easy

events Table:	
Column Name	Type
app_id	integer
event_type	string
timestamp	datetime

**Explanation:** Event\_type can be of either impression or click. We need to find the total number of impression and clicks and then use the formula

**Step1:** To get the number of clicks, create a new column 'clicks' filled as 1 whenever the event\_type=click and 0 if not. Similarly create another column 'impressions' using the same logic.

**Step2:** Now that we have the 2 fields for calculating the CTR, we can group by on the app\_id to find the sum of impressions and clicks for each app and calculate CTR using the formula. Filter for the required time period in the where clause.

```
WITH ctr_cte as
(
  SELECT
    app_id,
    CASE WHEN event_type = 'click' THEN 1 ELSE 0 END AS clicks,
    CASE WHEN event_type = 'impression' THEN 1 ELSE 0 END AS impressions
  FROM events
  WHERE timestamp >= '2022-01-01' AND timestamp < '2023-01-01'
```

)

```
select app_id, ROUND(100.0*sum(clicks)/sum(impressions),2) as ctr
from ctr_cte
GROUP BY app_id
```

---

22) Assume that you are given the table below containing information on various orders made by eBay customers. Write a query to obtain the user IDs and number of products purchased by the top 3 customers; these customers must have spent at least \$1,000 in total.

Output the user id and number of products in descending order. To break ties (i.e., if 2 customers both bought 10 products), the user who spent more should take precedence.

Company: Ebay

Difficulty: Easy

user_transactions Table:	
Column Name	Type
transaction_id	integer
product_id	integer
user_id	integer
spend	decimal

user_transactions Example Input:			
transaction_id	product_id	user_id	spend
131432	1324	128	699.78
131433	1313	128	501.00
153853	2134	102	1001.20
247826	8476	133	1051.00
247265	3255	133	1474.00
136495	3677	133	247.56

**Explanation:**

**Step1:** Find the count of products grouping by the customer\_id

**Step2:** Since the minimum overall purchase amount of the customers should be 1000, filter those customers using the having clause.

**Step3:** As only the top 3 customers are required we can use limit 3.

```
select user_id,
       count(product_id) as product_num
from user_transactions
group by user_id having sum(spend)>=1000
order by count(product_id) desc, sum(spend) desc limit 3
```

---

23) The LinkedIn Creator team is looking for power creators who use their personal profile as a company or influencer page. If someone's LinkedIn page has more followers than the company they work for, we can safely assume that person is a power creator.

Write a query to return the IDs of these LinkedIn power creators ordered by the IDs.

Assumption:

- Each person with a LinkedIn profile in this database works at one company only.

Company: LinkedIn

Difficulty: Easy

personal_profiles Table:	
Column Name	Type
profile_id	integer
name	string
followers	integer
employer_id	integer

personal_profiles Example Input:			
profile_id	name	followers	employer_id
1	Nick Singh	92,000	4
2	Zach Wilson	199,000	2
3	Daliana Liu	171,000	1
4	Ravit Jain	107,000	3
5	Vin Vashishta	139,000	6
6	Susan Wojcicki	39,000	5

**company\_pages Table:**

Column Name	Type
company_id	integer
name	string
followers	integer

**company\_pages Example Input:**

company_id	name	followers
1	The Data Science Podcast	8,000
2	Airbnb	700,000
3	The Ravit Show	6,000
4	DataLemur	200
5	YouTube	1,6000,000
6	DataScience.Vin	4,500

**Explanation:**

**Step1:** Join the personal\_profile and the company\_pages table using the employer\_id and the company\_id columns.

**Step2:** Now we can simply filter if the employee has more followers compared to the company

```
SELECT p.profile_id
FROM personal_profiles p JOIN company_pages c
      ON p.employer_id=c.company_id
WHERE p.followers > c.followers ORDER BY profile_id;
```

---

24) Your team at Accenture is helping a Fortune 500 client revamp their compensation and benefits program. The first step in this analysis is to manually review employees who are potentially overpaid or underpaid.

An employee is considered to be potentially overpaid if they earn more than 2 times the average salary for people with the same title. Similarly, an employee might be underpaid if they earn less than half of the average for their title. We'll refer to employees who are both underpaid and overpaid as compensation outliers for the purposes of this problem.

Write a query that shows the following data for each compensation outlier: employee ID, salary, and whether they are potentially overpaid or potentially underpaid (refer to Example Output below).

Company: Accenture

Difficulty: Medium

employee_pay Table:	
Column Name	Type
employee_id	integer
salary	integer
title	varchar

employee_pay Example Input:		
employee_id	salary	title
101	80000	Data Analyst
102	90000	Data Analyst
103	100000	Data Analyst
104	30000	Data Analyst
105	120000	Data Scientist

### **Explanation:**

**Step1:** Since each employee's salary has to be compared with avg.salary for the same job title, it is ideal to use the window functions to find the avg salary partitioned by job title.

**Step2:** Now that we have the employee salary and the average salary in each row. We can do a simple comparison to find if they are 'underpaid'(salary/2 < avg.salary) or 'overpaid'(salary\*2 > avg.salary) and assign it as 2 categories using the case statement. If the employees do not fall into either category then we can assign them as 'skip'

**Step3:** Filter out all the employees who are categorized as 'skip', since we want only underpaid and overpaid employees

```
select * from (
    SELECT employee_id, salary,
        CASE
            WHEN (salary*2) < avg(salary) OVER(PARTITION by title) THEN 'Underpaid'
            WHEN (salary/2) > avg(salary) OVER(PARTITION by title) THEN 'Overpaid'
            ELSE 'skip'
        END as status
    FROM employee_pay
) a
where status <> 'skip'
```

---

25) When you log in to your retailer client's database, you notice that their product catalog data is full of gaps in the category column. Can you write a SQL query that returns the product catalog with the missing data filled in?

### **Assumptions**

- Each category is mentioned only once in a category column.

- All the products belonging to the same category are grouped together.
- The first product from a product group will always have a defined category.
  - Meaning that the first item from each category will not have a missing category value.

Company: Accenture

Difficulty: Medium

products Table:	
Column Name	Type
product_id	integer
category	varchar
name	varchar

products Example Input:		
product_id	category	name
1	Shoes	Sperry Boat Shoe
2		Adidas Stan Smith
3		Vans Authentic
4	Jeans	Levi 511
5		Wrangler Straight Fit
6	Shirts	Lacoste Classic Polo
7		Nautica Linen Shirt

Example Output:		
product_id	category	name
1	Shoes	Sperry Boat Shoe
2	Shoes	Adidas Stan Smith
3	Shoes	Vans Authentic
4	Jeans	Levi 511
5	Jeans	Wrangler Straight Fit
6	Shirts	Lacoste Classic Polo
7	Shirts	Nautica Linen Shirt

### Explanation:

As the data is to be filled in a sequence from top to bottom we have to find a method to group all the NULL values to a category starting from the top. i.e. in the example input screenshot above product\_ids 1 to 3 should be grouped by a common value. Similarly product\_ids 4 and 5 should be grouped by a common value and 6,7 as a group.

**Step1:** We can assign a group(cat\_id) by taking the count of the products in each category and order them by product\_ids using window functions (t1 in the query). As we know if we use an aggregate function with 'order by' it will create a running aggregation(here running count). Since we did not use 'partition by', the entire table will be considered as one partition and everytime there is a new category the count is incremented by 1. Note: Since NULL means absence of value cat\_id will be filled with the same value as the previous row. Sample Screenshot of t1 query output

product_id	category	name	cat_id
1	Shoes	Sperry Boat Shoe	1
2		Adidas Stan Smith	1
3		Vans Authentic	1
4	Jeans	Levi 511	2
5		Wrangler Straight Fit	2
6	Shirts	Lacoste Classic Polo	3
7		Nautica Linen Shirt	3
8	Jackets	Belstaff Blue Fit	4
9		Uniqlo Comfort Jacket	4

**Step2:** We can use the output of the t1 query, filter out the rows where the category value is NULL to get the category\_id for each category column. Sample Screenshot of t2 query output

Output	
cat_id	category
1	Shoes
2	Jeans
3	Shirts
4	Jackets

**Step3:** Join the cat\_id from t1 and t2 to get the final output

```
with t1 as( select
    product_id, category, name,
    count(category) over (order by product_id) as cat_id
from products)
, t2 as (
    select cat_id, category from t1
    where category is not null
)
select t1.product_id, t2.category, t1.name
from t1 left join t2 on t2.cat_id = t1.cat_id
```

---

26) You're given two tables on Spotify users' streaming data. songs\_history table contains the historical streaming data and songs\_weekly table contains the current week's streaming data.

Write a query to output the user id, song id, and cumulative count of song plays as of 4 August 2022 sorted in descending order.

Definitions:

- songs\_weekly table currently holds data from 1 August 2022 to 7 August 2022.
- songs\_history table currently holds data up to 31 July 2022. The output should include the historical data in this table.

Assumption:

- There may be a new user or song in the songs\_weekly table not present in the songs\_history table.

Company: Spotify

Difficulty: Medium

songs_history Table:	
Column Name	Type
history_id	integer
user_id	integer
song_id	integer
song_plays	integer

songs_history Example Input:			
history_id	user_id	song_id	song_plays
10011	777	1238	11
12452	695	4520	1

song\_plays : Refers to the historical count of streaming or song plays by the user.

songs_weekly Table:	
Column Name	Type
user_id	integer
song_id	integer
listen_time	datetime

songs_weekly Example Input:		
user_id	song_id	listen_time
777	1238	08/01/2022 12:00:00
695	4520	08/04/2022 08:00:00
125	9630	08/04/2022 16:00:00
695	9852	08/07/2022 12:00:00

### Explanation:

**Step1:** The songs\_weekly table contains only information from August 1 to 7, however we need only info up to 4th August. So we can filter out data using the date function in the where clause

**Step2:** Get the count of songs listened by each user by using group by between the dates 1 to 4th August in the songs\_weekly table.

**Step3:** Since we need the overall information about the historical songs listened by each users along with the songs listened in the current week, we can union them together. Note: The column names and order should be in the matching

**Step4:** Now we have all the information in one single table, we can simply take a sum of all the songs listened grouped by each user\_id

WITH filtered AS (

SELECT user\_id, song\_id, song\_plays FROM songs\_history



```

        UNION ALL
SELECT user_id, song_id, COUNT(song_id) AS song_plays FROM songs_weekly
WHERE DATE_PART('day', listen_time) < 5
GROUP BY user_id, song_id
)

```

```

SELECT user_id, song_id, SUM(song_plays) AS song_plays FROM filtered
GROUP BY user_id, song_id
ORDER BY song_plays DESC;

```

---

27) Assume you are given the table below containing the information on the searches attempted and the percentage of invalid searches by country. Write a query to obtain the percentage of invalid searches.

Output the country in ascending order, total searches and overall percentage of invalid searches rounded to 2 decimal places.

Notes:

- num\_search = Number of searches attempted; invalid\_result\_pct = Percentage of invalid searches.
- In cases where countries have search attempts but do not have a percentage of invalid searches in invalid\_result\_pct, it should be excluded, and vice versa.
- To find the percentages, multiply by 100.0 and not 100 to avoid integer division.

Difficulty: Medium

Company: Google

search_category Table:	
Column Name	Type
country	string
search_cat	string
num_search	integer
invalid_result_pct	decimal

search_category Example Input:			
country	search_cat	num_search	invalid_result_pct
UK	home	null	null
UK	tax	98000	1.00
UK	travel	100000	3.25

**Explanation:** The 'invalid\_result\_pct' is the percentage of invalid results for each of the search categories. However we need to find the overall invalid\_result\_pct which is not a simple addition. So we need to find the actual invalid search count first, add them up at country level and then find the percentage.

**Step1:** Find the actual invalid searches for each category by  $(\text{num\_search} * \text{invalid\_result\_pct}) / 100$  and remove the searches where the num\_searches is NULL in the where clause.

**Step2:** Now find the sum of overall search and also find the overall invalid search percentage by  $(\text{summing the overall invalid searches}) / (\text{sum of overall searches}) * 100.0$  by grouping by on the country

```
select country, sum(num_search) as total_search,
round((SUM(inval)*100.0/sum(num_search),2) as invalid_searches_pct
from
(
  SELECT *,
    (num_search*invalid_result_pct)/100 as inval
  FROM search_category
  where num_search is not null and invalid_result_pct is not null
) as c
group by country
ORDER BY country
```

---

28) For every customer that bought Photoshop, return a list of the customers, and the total spent on all the products except for Photoshop products.

Sort your answer by customer ids in ascending order

Difficulty: Medium

Company: Adobe

adobe_transactions Table:	
Column Name	Type
customer_id	integer
product	string
revenue	integer

adobe_transactions Example Input:		
customer_id	product	revenue
123	Photoshop	50
123	Premier Pro	100
123	After Effects	50
234	Illustrator	200
234	Premier Pro	100

### **Explanation**

**Step1:** First step, find all the customers who have purchased 'Photoshop' i.e. inner query

**Step2:** Filter only those customers who have bought the photoshop product from the previous step from the same table. Now remove all the rows which correspond to transactions related to Photoshop.

**Step3:** Group by customer\_id and find the sum of overall revenue which results in customer who have purchased photoshop and the overall revenue is from other products other than photoshop

```
SELECT customer_id,
       SUM(revenue) AS revenue
FROM adobe_transactions
WHERE customer_id IN
      (SELECT customer_id
       FROM adobe_transactions
       WHERE product = 'Photoshop'
      )
AND product <> 'Photoshop'
GROUP BY customer_id;
```

---

29) You are given a table of PayPal payments showing the payer, the recipient, and the amount paid. A two-way unique relationship is established when two people send money back and forth. Write a query to find the number of two-way unique relationships in this data.

Assumption:

A payer can send money to the same recipient multiple times.

Company: Paypal

Difficulty: Medium

payments Table:		payments Example Input:		
Column Name	Type	payer_id	recipient_id	amount
payer_id	integer	101	201	30
recipient_id	integer	201	101	10
amount	integer	101	301	20
		301	101	80
		201	301	70

**Explanation:**

Find the userA, userB combination. userA can be a sender and userB can be a receiver or vice versa. Both of these are considered as a single unique relationship between userA and userB..

**Step1:** To find if two users have been both a sender and receiver, self join the two tables by swapping the payer\_id and recipient\_id.

**Step2:** All the records that match are the ones which have a two way relationship.

**Step3:** To represent relationship between two users only once we filter receipt\_id > payer\_id (can also be receipt\_id < payer\_id). This will result in a two way relationship between users occurring only once.

**Step4:** Finally take a count to get the unique users relationship combination

```
select count(*) as comb
from(
    SELECT
        distinct a.payer_id,
        a.recipient_id
    FROM payments a join payments b
        on a.payer_id = b.recipient_id and
        b.payer_id = a.recipient_id
    Where a.amount is not null and b.amount is not null
    and a.payer_id < a.recipient_id
) x
```

---

30) Assume you are given the table below containing information on user purchases. Write a query to obtain the number of users who purchased the same product on two or more different days. Output the number of unique users.

Company: Stitch fix

Difficulty: Medium

purchases Table:		purchases Example Input:			
Column Name	Type	user_id	product_id	quantity	purchase_date
user_id	integer	536	3223	6	01/11/2022 12:33:44
product_id	integer	827	3585	35	02/20/2022 14:05:26
quantity	integer	536	3223	5	03/02/2022 09:33:28
purchase_date	datetime	536	1435	10	03/02/2022 08:40:00
		827	2452	45	04/09/2022 00:00:00

**Explanation:**

**Step1:** To find if a customer has purchased the same product for more than 1 day, then there is going to be more than one transaction record for that product for that user with different dates.

**Step2:** With this logic we can group by user\_id, product\_id and find the count of unique purchase dates. If the purchase\_dates are more than 1 which means the customer has purchased that product for more than 2 days.

**Step3:** Only select those unique customers who have had any purchase where the purchase date Count >1.

```
with cte as (  
    Select user_id,  
           product_id,  
           count(DISTINCT date(purchase_date)) as rep_purchase_count  
    from purchases  
    group by user_id, product_id  
)  
  
select count(distinct (user_id)) as repeat_purchasers  
from cte where rep_purchase_count>1
```

---